

# Dynamic Programming: Edit Distance

## Programming Assignment

---

In this assignment you will compute DNA sequence alignments. The optimization process involves evaluating how well individual base pairs match up in the DNA sequence case.

## Global Sequence Alignment

---

Write a program to compute the optimal sequence alignment of two DNA strings. This program will introduce you to the emerging field of *computational biology* in which computers are used to do research on biological systems. Further, you will be introduced to a powerful algorithmic design paradigm known as *dynamic programming*.

**Biology review.** A *genetic sequence* is a string formed from a four-letter alphabet {Adenine (A), Thymine (T), Guanine (G), Cytosine (C)} of biological macromolecules referred to together as the DNA bases. A *gene* is a genetic sequence that contains the information needed to construct a protein. All of your genes taken together are referred to as the human genome, a blueprint for the parts needed to construct the proteins that form your cells. Each new cell produced by your body receives a copy of the genome. This copying process, as well as natural wear and tear, introduces a small number of changes into the sequences of many genes. Among the most common changes are the substitution of one base for another and the deletion of a substring of bases; such changes are generally referred to as *point mutations*. As a result of these point mutations, the same gene sequenced from closely related organisms will have slight differences.

**The problem.** Through your research you have found the following sequence of a gene in a previously unstudied organism.

A A C A G T T A C C

What is the function of the protein that this gene encodes? You could begin a series of uninformed experiments in the lab to determine what role this gene plays. However, there is a good chance that it is a variant of a known gene in a previously studied organism. Since biologists and computer scientists have laboriously determined (and published) the genetic sequence of many organisms (including humans), you would like to leverage this information to your advantage. We'll compare the above genetic sequence with one which has already been sequenced and whose function is well understood.

T A A G G T C A

If the two genetic sequences are similar enough, we might expect them to have similar functions. We would like a way to quantify "similar enough."

**Edit-distance.** In this assignment we will measure the similarity of two genetic sequences by their *edit distance*, a concept first introduced in the context of coding theory, but which is now widely used in spell checking, speech recognition, plagiarism detection, file revisioning, and computational linguistics. We align the two sequences, but we are permitted to *insert gaps* in either sequence (e.g., to make them have the same length). We pay a penalty for each gap that we insert and also for each pair of characters that mismatch in the final alignment. Intuitively, these penalties model the relative likeliness of point mutations arising from deletion/insertion and substitution. We produce a numerical score according to the following table, which is widely used in biological applications:

<i>operation</i>	<i>cost</i>
<i>insert a gap</i>	2
<i>align two characters that mismatch</i>	1
<i>align two characters that match</i>	0

Here are two possible alignments of the strings  $x = \text{"AACAGTTACC"}$  and  $y = \text{"TAAGGTCA"}$ :

x	y	cost	x	y	cost
A	T	1	A	T	1
A	A	0	A	A	0
C	A	1	C	-	2
A	G	1	A	A	0
G	G	0	G	G	0
T	T	0	T	G	1
T	C	1	T	T	0
A	A	0	A	-	2
C	-	2	C	C	0
C	-	2	C	A	1
---			---		
8			7		

The first alignment has a score of 8, while the second one has a score of 7. The *edit-distance* is the score of the best possible alignment between the two genetic sequences over all possible alignments. In this example, the second alignment is in fact optimal, so the edit-distance between the two strings is 7. Computing the edit-distance is a nontrivial computational problem because we must find the best alignment among exponentially many possibilities.

We will explain a recursive solution which is an elegant approach. However it is far too inefficient because it recalculates each subproblem over and over. Once we have defined the recursive definition we can redefine the solution using a dynamic programming approach which calculates each subproblem once.

**A recursive solution.** We will calculate the edit-distance between the two original strings  $x$  and  $y$  by solving many edit-distance problems on the *suffixes* of the two strings. We use the notation  $x[i]$  to refer to character  $i$  of the string. We also use the notation  $x[i..M]$  to refer to the suffix of  $x$  consisting of the characters  $x[i], x[i+1], \dots, x[M-1]$ . Finally, we use the notation  $\text{opt}[i][j].\text{cost}$  to denote the edit distance of  $x[i..M]$  and  $y[j..N]$ . For example, consider the two strings  $x = \text{"AACAGTTACC"}$  and  $y = \text{"TAAGGTCA"}$  of length  $M = 10$  and  $N = 8$ , respectively. Then,  $x[2]$  is 'C',  $x[2..M]$  is "CAGTTACC", and  $y[8..N]$  is the empty string. The edit distance of  $x$  and  $y$  is  $\text{opt}[0][0].\text{cost}$ .

Now we describe a recursive scheme for computing the edit distance of  $x[i..M]$  and  $y[j..N]$ . Consider the first pair of characters in an optimal alignment of  $x[i..M]$  with  $y[j..N]$ . There are three possibilities:

1. The optimal alignment matches  $x[i]$  up with  $y[j]$ . In this case, we pay a penalty of either 0 or 1, depending on whether  $x[i]$  equals  $y[j]$ , plus we still need to align  $x[i+1..M]$  with  $y[j+1..N]$ . What is the best way to do this? This subproblem is exactly the same as the original sequence alignment problem, except that the two inputs are each suffixes of the original inputs. Using our notation, this quantity is  $\text{opt}[i+1][j+1].\text{cost}$ .
2. The optimal alignment matches the  $x[i]$  up with a gap. In this case, we pay a penalty of 2 for a gap and still need to align  $x[i+1..M]$  with  $y[j..N]$ . This subproblem is identical to the original sequence alignment problem, except that the first input is a proper suffix of the original input.
3. The optimal alignment matches the  $y[j]$  up with a gap. In this case, we pay a penalty of 2 for a gap and still need to align  $x[i..M]$  with  $y[j+1..N]$ . This subproblem is identical to the original sequence alignment problem, except that the second input is a proper suffix of the original input.

The key observation is that all of the resulting subproblems are sequence alignment problems on suffixes of the original inputs. To summarize, we can compute

$\text{opt}[i][j].\text{cost}$  by taking the minimum of three quantities:

$$\text{opt}[i][j].\text{cost} = \min \{ \text{opt}[i+1][j+1].\text{cost} + 0/1, \text{opt}[i+1][j].\text{cost} + 2, \text{opt}[i][j+1].\text{cost} + 2 \}$$

This equation works assuming  $i < M$  and  $j < N$ . Aligning an empty string with another string of length  $k$  requires inserting  $k$  gaps. Thus, in general we should set

$\text{opt}[M][j].\text{cost} = 2 + \text{opt}[M][j+1].\text{cost}$  and  $\text{opt}[i][N].\text{cost} = 2 + \text{opt}[i+1][N].\text{cost}$ . The simplest way to handle this situation is to add an extra row and column to the end of the  $\text{opt}$  matrix that corresponds to matching a gap at the end of each string. For our example, the final matrix is:

			0	1	2	3	4	5	6	7	8
x\y			T	A	A	G	G	T	C	A	-
0	A		7	8	10	12	13	15	16	18	20
1	A		6	6	8	10	11	13	14	16	18
2	C		6	5	6	8	9	11	12	14	16
3	A		7	5	4	6	7	9	11	12	14
4	G		9	7	5	4	5	7	9	10	12
5	T		8	8	6	4	4	5	7	8	10
6	T		9	8	7	5	3	3	5	6	8
7	A		11	9	7	6	4	2	3	4	6
8	C		13	11	9	7	5	3	1	3	4
9	C		14	12	10	8	6	4	2	1	2
10	-		16	14	12	10	8	6	4	2	0

By examining `opt[0][0].cost`, we conclude that the edit distance of `x` and `y` is 7.

**A dynamic programming approach.** A direct implementation of the above recursive scheme will work, but it is spectacularly inefficient. If both input strings have `N` characters, then the number of recursive calls will exceed  $2^N$ . To overcome this performance bug, we use *dynamic programming*. Dynamic programming is a powerful algorithmic paradigm, first introduced by Bellman in the context of operations research, and then applied to the alignment of biological sequences by Needleman and Wunsch. Dynamic programming now plays the leading role in many computational problems, including control theory, financial engineering, and bioinformatics, including **BLAST** (the sequence alignment program almost universally used by molecular biologist in their experimental work). **The key idea of dynamic programming is to break up a large computational problem into smaller subproblems, store the answers to those smaller subproblems, and, eventually, use the stored answers to solve the original problem.** This avoids recomputing the same quantity over and over again. Instead of using recursion, use a nested loop that calculates `opt[i][j].cost` in the *right* order so that `opt[i+1][j+1].cost`, `opt[i+1][j].cost`, and `opt[i][j+1].cost` are all computed before we try to compute `opt[i][j].cost`.

**Recovering the alignment itself.** The above procedure describes how to compute the edit distance between two strings. We now outline how to recover the optimal alignment itself. The key idea is to retrace the steps of the dynamic programming algorithm backwards, re-discovering the path of choices (highlighted in red in the table above) from `opt[0][0].cost` to `opt[M][N].cost`. To determine the choice that led to `opt[i][j].cost`, we consider the three possibilities:

1. The optimal alignment matches `x[i]` up with `y[j]`. In this case, we must have `opt[i][j].cost = opt[i+1][j+1].cost` if `x[i]` equals `y[j]`, or `opt[i][j].cost = opt[i+1][j+1].cost + 1` otherwise.
2. The optimal alignment matches `x[i]` up with a gap. In this case, we must have `opt[i][j].cost = opt[i+1][j].cost + 2`.

3. The optimal alignment matches  $y[j]$  up with a gap. In this case, we must have  $\text{opt}[i][j].\text{cost} = \text{opt}[i][j+1].\text{cost} + 2$ .

Depending on which of the three cases apply, we move diagonally, down, or right towards  $\text{opt}[M][N].\text{cost}$ , printing out  $x[i]$  aligned with  $y[j]$  (case 1),  $x[i]$  aligned with a gap (case 2), or  $y[j]$  aligned with a gap (case 3). In the example above, we know that the first T aligns with the first A because  $\text{opt}[0][0].\text{cost} = \text{opt}[1][1].\text{cost} + 1$ , but  $\text{opt}[0][0].\text{cost} \neq \text{opt}[1][0].\text{cost} + 2$  and  $\text{opt}[0][0].\text{cost} \neq \text{opt}[0][1].\text{cost} + 2$ . The optimal alignment is:

x	y	cost
-----		
A	T	1
A	A	0
C	-	2
A	A	0
G	G	0
T	G	1
T	T	0
A	-	2
C	C	0
C	A	1

**Your program.** Write a program `EditDistance.cpp` whose `main` method reads, from standard input, two strings of characters. (Although, in the application described, the characters represent genetic sequences, your program should handle any sequence of alphanumeric characters.) Your program should then print the edit distance between the two strings, and print out the optimal match along with the individual penalties using the following format:

- The first line should contain the edit distance, preceded by the text "Edit distance = ".
- Each subsequent line should contain a character from the first string, followed by the paired character from the second string, followed by the associated penalty. Use the character '-' to indicate a gap in either string. To handle the base case of aligning strings, you will have to allow the last  $k$  characters of one string to match gaps in the other string. In particular, the final node in your path will match a "gap" at the end of the first string to a "gap" at the end of the second string, with a penalty of 0. You should *not* print this final node of your path.

Here is a sample output:

```
Edit distance = 7
A T 1
A A 0
C - 2
A A 0
G G 0
T G 1
T T 0
A - 2
C C 0
C A 1
```

**Testing** The data sequence.zip contains short test data files and actual genomic data files. In addition, you should create test cases of your own to make sure you handle all special cases correctly.

**Submission.** Submit all files in a single zip file. Do not include the test cases you downloaded with the assignment. Do include the test cases you created and they should all be the root folder. Your report should specify what additional test data you created to test your program for special cases? (At least one set of test data required) and how did you test for correctness?

---

*This assignment was taken from upenn.*